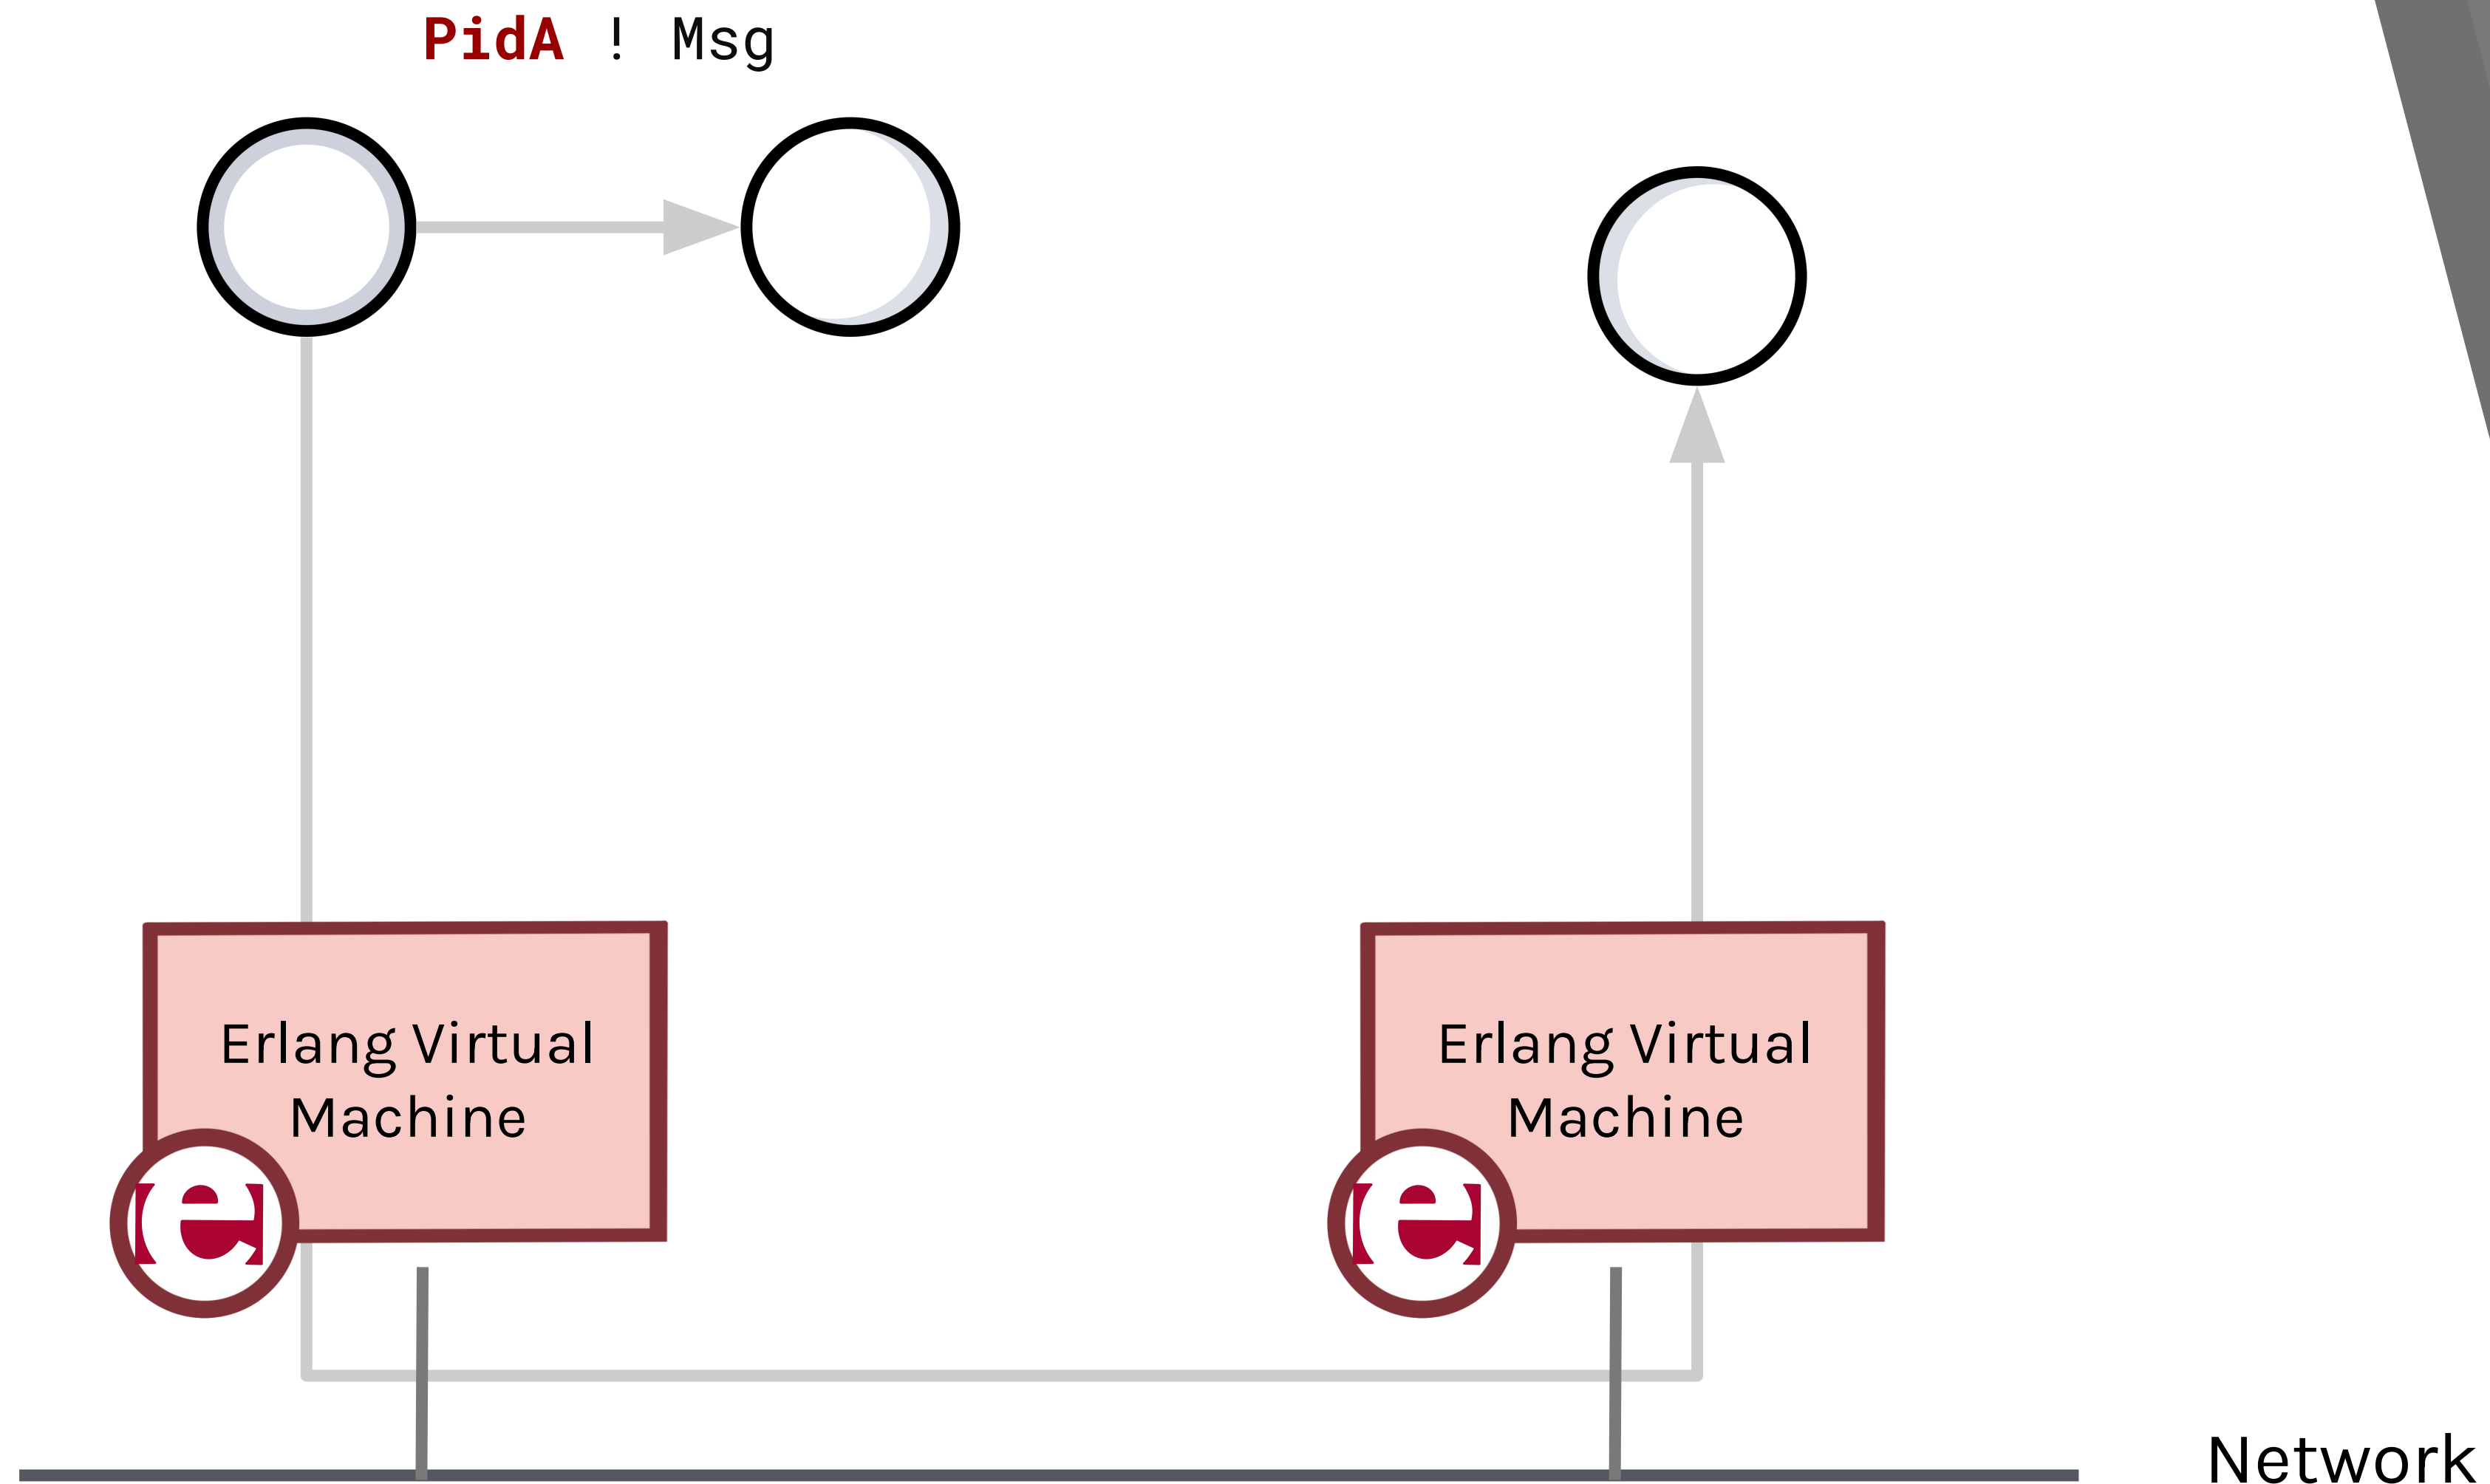


# Distributed Programming


# Distributed Programming

- ▶ Distributed Systems
- ▶ Erlang Distribution
- ▶ Node Connections
- ▶ Distributed BIFs
- ▶ Net Kernel
- ▶ Hidden Nodes
- ▶ Firewalls and Erlang Distribution

# Distributed Systems



# Distributed Systems



- ▶ A distributed system consists of a program running on a collection of loosely coupled processors

# Distributed Systems

- ▶ **There are many reasons for building distributed systems**
  - Performance
  - Reliability and fault tolerance
  - Accessing remote resources
  - Distributed applications (possibly heterogeneous)
  - Extendibility and Scalability

# Erlang Distribution: **example**

```
$ erl -sname foo
...
(foo@node)1> S = self().
<0.37.0>
(foo@node)2> spawn('bar@node', s, send_node, [S]).
<5827.42.0>
(foo@node)3> flush().
Shell got 'bar@node'
```

**s:send\_node(S)** runs on **bar@node**

```
$ erl -sname bar
...
(bar@node)1>
```

```
-module(s).
-export([send_node/1]).

send_node(To) -> To ! node().
```

# Erlang Distribution

- ▶ A **node** is an executing instance of the Erlang VM
- ▶ A node is said to be **alive** if it can communicate with other nodes
- ▶ Nodes that are alive have a unique name on that host
- ▶ **Name@Host** is the identifier of an Erlang node
- ▶ A distributed Erlang node is started using the **-name** and **-sname** flag
  - In Windows, use the DOS/Command prompt to run the command



# Erlang Distribution

- ▶ Erlang nodes can communicate only if they share the same cookie or if they know each other's cookies
- ▶ Simple on trusted networks and local clusters
- ▶ Harder on untrusted networks



# Erlang Distribution: **starting**

```
cesarini@caen> erl -sname foo -setcookie hello
```

*will start a node with the short name foo@caen*

```
cesarini@caen> erl -name foo
```

*will start a node with the long name foo@caen.ericsson.se*

- ▶ Nodes started with long names may not communicate with nodes using short names
- ▶ Nodes using short names do not require a DNS server
- ▶ If no cookie is defined when starting up the system, then one will be created and stored in the **~/.erlang.cookie** file

# Messages: **example**

```
(foo@node)4> spawn('bar@node', s, s, []).  
<5827.47.0>  
(foo@node)5> {server, 'bar@node'} ! {hi, self()}.  
{hi, <0.37.0>}  
(foo@node)6> flush().  
Shell got hi
```

**hi** sent from **foo@node** to **bar@node**, which replies with **hi**

```
-module(s).  
-export([s/0]).  
  
s() -> register(server, self()), l().  
  
l() -> receive {M, Pid} -> Pid ! M end, l().
```

# Erlang Distribution: **messages**

```
Pid ! Message  
{Name, Node} ! Msg
```

- ▶ Send messages to processes on remote nodes transparently using the **!** construct
- ▶ Messages will be delivered in the same order they are sent.
- ▶ Only difference is that the remote node might go down.
- ▶ With little changes (or doing it right from the start), programs running on a single processor can easily become distributed

# Distributed Example

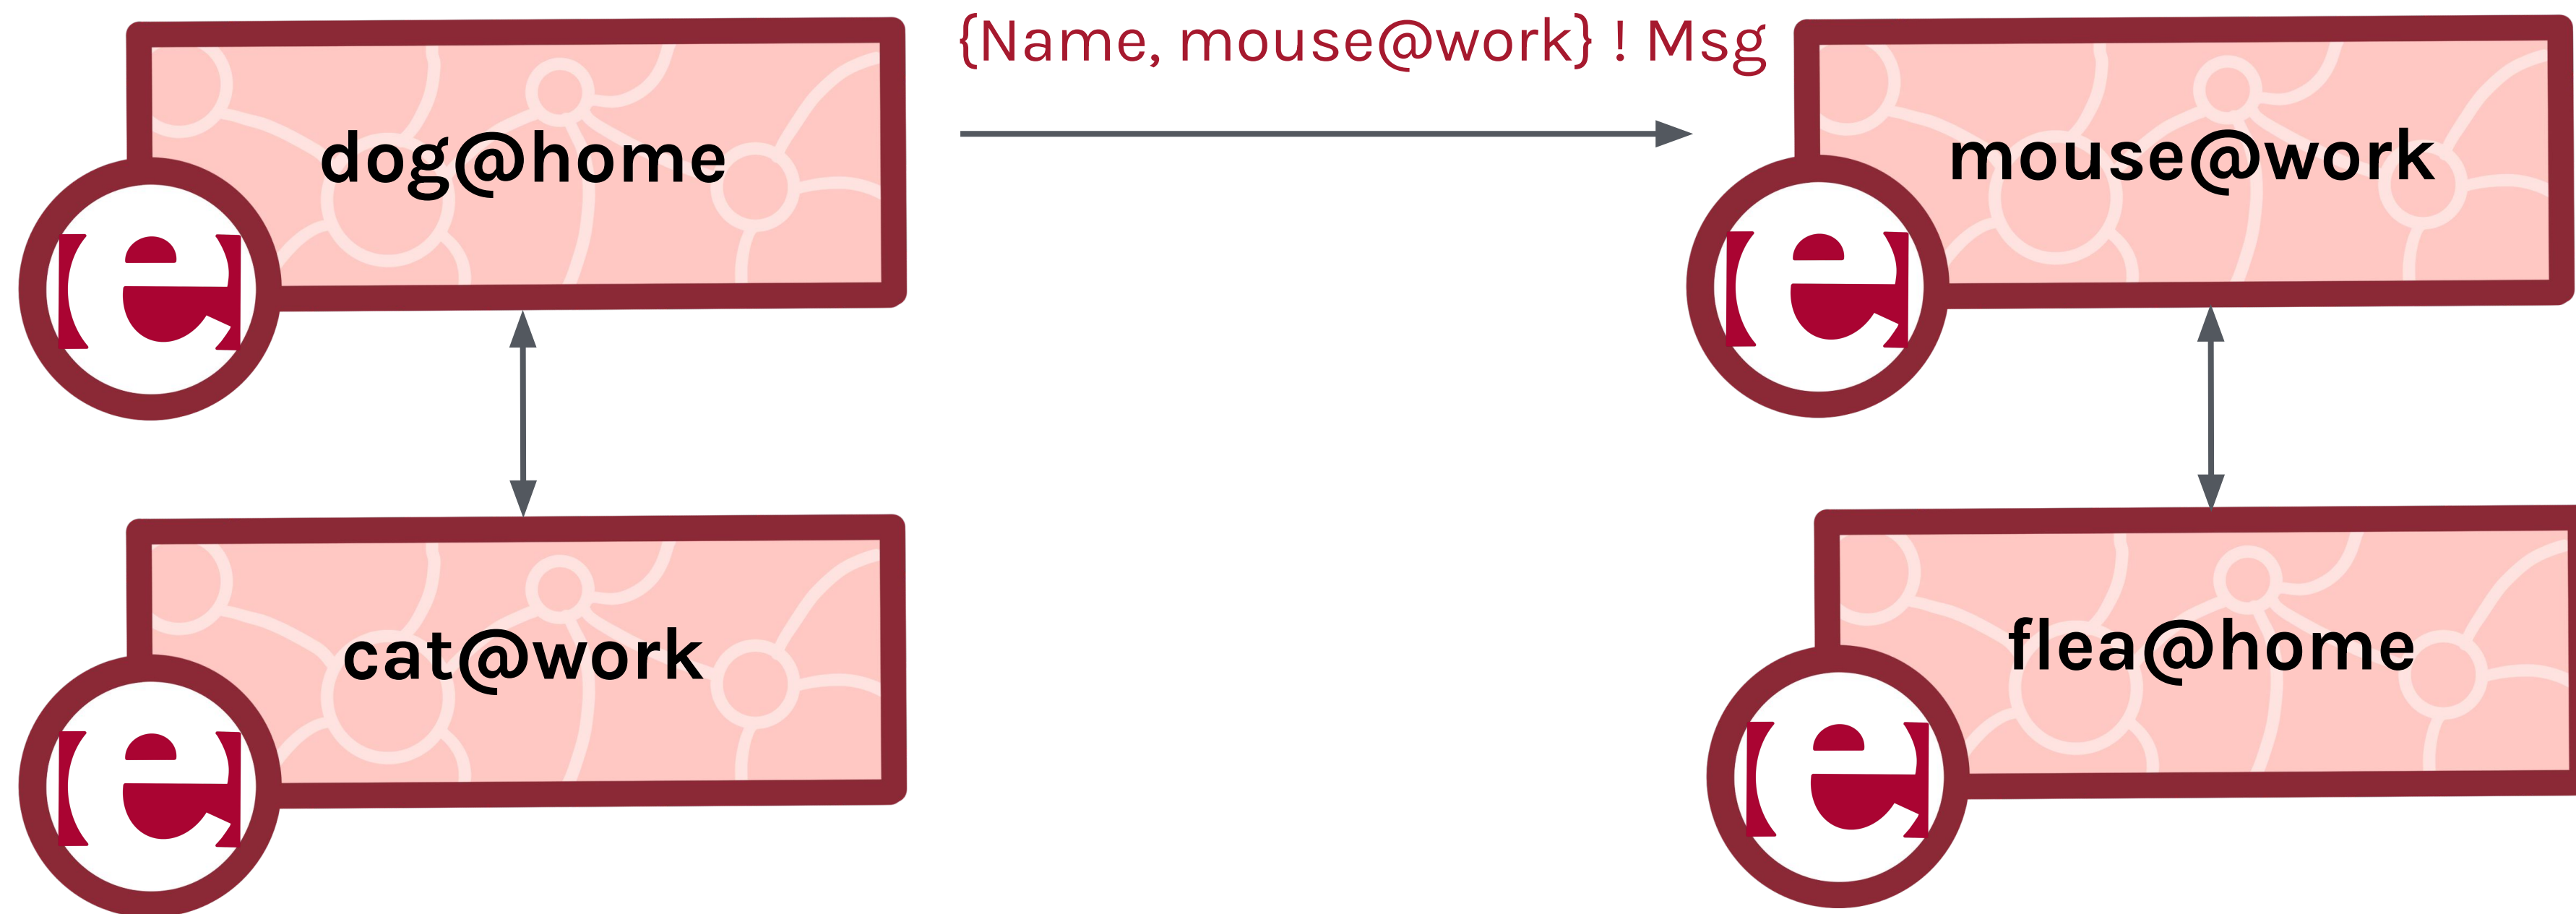
```
call(Message, ServerNode) ->  
    {frequency, ServerNode} ! {request, self(), Message},  
    receive  
        {reply, Reply} -> Reply  
    end.  
reply(Pid, Message) ->  
    Pid ! {reply, Message}.
```

- ▶ Recall the resource server call? We can now receive requests to safely allocate resources from other nodes

# Node Connections

- ▶ Nodes are loosely connected to each other
  - Connections are based on TCP/IP
- ▶ Connections are explicitly set up by the run time system when a node is first referred to
  - They are not set up by the programmer
- ▶ Connected nodes share information on new node
- ▶ Nodes monitor each other checking if their peers are alive
- ▶ Nodes can come and go dynamically in a similar manner as processes


# Node Connections



- ▶ dog@home is connected to cat@work and mouse@work to flea@home
- ▶ A process in dog@home sends a message to a registered process in mouse@work



# Node Connections



- ▶ dog@home sets up a connection to mouse@work and is informed of flea@home
- ▶ dog@home tells cat@work of mouse@work and flea@home, and tells mouse@work of cat@work
- ▶ This goes on and they all get happily connected



# Distributed BIFs

- ▶ `Pid = spawn(Node, Mod, Func, Args)`
- ▶ `Pid = spawn_link(Node, Mod, Func, Args)`
- ▶ `monitor_node(Node, Flag)`
- ▶ `MyName@Host = node()`
- ▶ `[Node|Nodes] = nodes()`
- ▶ `Node = node(Item) % Item is a PID, Ref or Port`
- ▶ `disconnect_node(Node)`
- ▶ `erlang:set_cookie(Node, Cookie)`

# Net Kernel

- ▶ The net\_kernel is an Erlang process that coordinates operations in a distributed node
- ▶ BIFs such as **spawn/4** are converted by net\_kernel to messages, and sent to the net\_kernel on the remote node
- ▶ Handles authentication and rejects bad cookies
- ▶ You can change the net\_kernel with a user-defined process giving a specific behaviour
  - Changing the authentication scheme,
  - not allowing remote nodes to spawn processes, etc

# Hidden Nodes: **example**

```
erl -sname a
...

(a@node)1>
  net_kernel:connect('b@node'),
  net_kernel:connect('c@node').
true
(a@node)2> nodes().
['b@node']
(a@node)3> nodes(hidden).
['c@node']
(a@node)4> nodes(connected).
['b@node', 'c@node']
```

```
erl -sname b
...

(b@node)1> nodes().
['a@node']
(b@node)2> nodes(connected)
['a@node']
```

```
erl -sname c -hidden
...

(c@node)1> nodes().
[]
(c@node)2> nodes(hidden).
['a@node']
```

# Hidden Nodes

```
erl -sname foo -hidden  
will start a hidden node
```

- ▶ Hidden nodes have to be connected individually
- ▶ Typically used for Operation and Maintenance
- ▶ Can be used to create hierarchies of nodes

# Firewalls and Erlang Distribution

- ▶ Poke a hole in the firewall, port 4369
  - Erlang Port Mapper Daemon (EPMD), listens to incoming connection requests
- ▶ Poke a hole in the firewall, ports 9100 to 9105
  - Set undocumented kernel variables:
    - `application:set_env(kernel, inet_dist_listen_min, 9100)`
    - `application:set_env(kernel_inet_dist_listen_max, 9105)`
    - Forces Erlang to use ports 9100 to 9105 for distribution
- ▶ Tunnel over using SSH
- ▶ Roll your own distribution handler

# Erlang Port Mapper Daemon (EPMD)

- ▶ The EPMD is a name server on hosts involved in distributed Erlang
- ▶ It is used to find IP address and port of remote hosts
- ▶ Only the local node name is stored, not the host
- ▶ Can be used interactively to query:
  - names of currently running nodes
  - remove nodes from currently known nodes

# Distributed Programming

- ▶ Distributed Systems
- ▶ Erlang Distribution
- ▶ Node Connections
- ▶ Distributed BIFs
- ▶ Net Kernel
- ▶ Hidden Nodes
- ▶ Firewalls and Erlang Distribution