

Contents

1 Starting the System and Basic Erlang	3
The Shell	3
Setting up an editor	4
IntelliJ IDEA	4
Atom	4
Emacs	4
Modules and Functions	5
Temperature Conversions	5
Simple Pattern Matching	6
2 Sequential Programming	7
Evaluating Expressions	7
Creating Lists	7
Side Effects	7
Database Handling Using Lists	8
ADVANCED: Manipulating Lists	9
ADVANCED: Database Handling using Trees	10
3 Concurrent Programming	11
An Echo Server	11
The Process Ring	11
ADVANCED: The Process Crossring	12
4 Process Design Patterns	14
A Database Server	14
A Turnstile	14
ADVANCED: A Database Server with transactions	15
5 Process Error Handling	16
The Linked Ping Pong Server	16
A Reliable Mutex Semaphore	16
A Reliable Database Server with transactions	16
ADVANCED: A Supervisor Process	17
6 Functional Programming Constructs	19
Higher Order Functions	19
List Comprehensions	20
7 Records and Maps	21
Database Handling Using Records	21
Database Handling Using Maps	21
8 Erlang Term Storage	22
Database handling using ETS	22

1 Starting the System and Basic Erlang

These exercises will help you get accustomed with the Erlang development and run time environments. The exercises marked Advanced are harder or more complex exercises to push your understanding of Erlang further. Once you have set up the Erlang mode for Emacs, you will be ready to write your first program.

To start the Erlang shell, type `erl` when working on Unix environments or double click on the Erlang Icon in Windows environments. Once your shell has started, you will get some system printouts followed by a prompt.

If you are working in Unix, you should get something like this (Possibly with more system printouts).

Unix Shell
<code>\$ erl</code>

Erlang Shell Session
<code>Eshell V10.5 (abort with ^G) 1></code>

The Shell

Type in the following Erlang expressions in the shell. They will show some of the principles (including pattern matching and single variable assignment) described in the lectures. What happens when they execute, what values do the expressions return, and why?

A. Erlang expressions

```
1 + 1.  
[1|[2|[3|[]]]].
```

B. Assigning through pattern matching

```
A = 1.  
B = 2.  
A + B.  
A = A + 1.
```

C. Recursive lists definitions

```
L = [A|[2,3]].  
[[3,2]|1].  
[H|T] = L.
```

D. Flow of execution through pattern matching

```
B = 2.  
B = 2.  
2 = B.  
B = C.  
C = B.  
B = C. (repeat it now that C is bound).
```

E. Extracting values in composite data types through pattern matching

```
Person = {person, "Mike", "Williams", [1,2,3,4]}.  
{person, Name, Surname, Phone} = Person.  
Name.
```

Setting up an editor

IntelliJ IDEA

A free and easy to learn IDE is offered by JetBrains, with decent Erlang and Elixir support via plugins. Windows, Mac OS X and Linux all are supported.

After installation, proceed to the Options ⇒ Plugins ⇒ Browse Repositories, and there find Erlang or Elixir.

<https://www.jetbrains.com/idea/download/>

Atom

An editor with an easy to learn interface, large amount of convenience plugins and multiple languages supported. Supports both Erlang and Elixir. Windows, Mac OS X and Linux all are supported.

After installation, proceed to the Options ⇒ Plugins, and there find Erlang or Elixir.

<https://atom.io/>

Emacs

An Erlang mode for Emacs exists. Detailed instructions on how to setup the erlang-mode on a UNIX or Windows machine are available at:

http://www.erlang.org/doc/apps/tools/erlang_mode_chapter.html

You can also install the erlang-mode for Emacs using a MELPA package:

<http://melpa.milkbox.net/#/erlang>

The standard erlang-mode for Emacs provides basic functionalities such as syntax highlighting, indentation and in-module code navigation. For more advanced features you may want to look at the Erlang Development Tool Suite (EDTS):

<https://github.com/tjarvstrand/edts>

Modules and Functions

Copy the demo module from the *Modules* example slide in the *Basic Erlang* course material. Compile it and try and run it from the shell. What happens when you call `demo:times(3, 5)`? What about `double(6)` when omitting the module name?

Erlang Shell Session
Eshell V10.5 (abort with ^ G) 1> c(demo). {ok, demo} 2> demo:double(6). 12

Temperature Conversions

Part A

Write functions `temp:f2c(Fahrenheit)` and `temp:c2f(Celsius)` which convert between Fahrenheit and Celsius temperature scales in a module named *temp.erl*.

Hint:

$$5 * (F - 32) = 9 * C$$

Part B

Write a function `temp:convert(Temperature)` which combines the functionality of `f2c` and `c2f`. A usage example follows:

Erlang Shell Session
Eshell V10.5 (abort with ^ G) 1> temp:convert({c, 100}). {f, 212.0}. 2> temp:convert({f, 32}). {c, 0.0}.

Simple Pattern Matching

Boolean operators compare values that are either true or false and also return true or false, depending on the input. As an example, the and operator will return true if, and only if both operands are true; otherwise, false will be returned. Another operator is or, which will return true if any of the operands is true:

P	Q	P and Q	P or Q
false	false	false	false
false	true	false	true
true	false	false	true
true	true	true	true

Table 1: Boolean Operators

A third operator, called not, will simply reverse the input it has received: not true will return false and not false will return true.

Write a module `boolean.erl` that takes logical expressions and boolean values (represented as the atoms `true` and `false`) and returns their boolean result. The functions you should write should include `b_not/1`, `b_and/2` and `b_or/2`. You may not use the logical constructs `and`, `or` or `not`. Test your module from the shell.

```
b_not(false) -> true
b_and(false, true) -> false
b_and(b_not(b_and(true, false)), true) -> true
```

Note:

`foo(X) -> Y` means that calling the function `foo` with the parameter `X` will result in the value `Y` being returned.

`and`, `or` and `not` are reserved words in Erlang.

2 Sequential Programming

These exercises will get you familiar with recursion and its different uses. Pay special attention to the different recursive patterns that we covered during the lectures. If you are having problems finding bugs or following the recursion, try using the debugger.

Evaluating Expressions

Part A

Write a function `sum/1` which given a positive integer N will return the sum of all the integers between 1 and N .

```
sum(5) -> 15.
```

Part B

Write a function `sum_interval/2` which given two integers N and M , where $N \leq M$, will return the sum of the interval between N and M . If $N > M$, you want your process to terminate abnormally.

```
sum_interval(1,3) -> 6.
```

```
sum_interval(6,6) -> 6.
```

Creating Lists

Part A

Write a function `create/1` which returns a list of the format $[1, 2, \dots, N-1, N]$.

```
create(3) -> [1,2,3].
```

Part B

Write a function `reverse_create/1` which returns a list of the format $[N, N-1, \dots, 2, 1]$.

```
reverse_create(3) -> [3,2,1].
```

Side Effects

Part A

Write a function `print/1` which prints out the integers between 1 and N .

Erlang Shell Session
Eshell V10.5 (abort with ^ G) 1> print(5). 1

```
2
3
4
5
ok
```

Hint:

```
Use io:format("Number:~p~n", [N]).
```

Part B

Write a function `even_print/1` which prints out the even integers between `1` and `N`.

Erlang Shell Session

```
Eshell V10.5 (abort with ^ G)

1> even_print(5).
2
4
ok
```

Database Handling Using Lists

Write a module `db.erl` that creates a database and is able to store, retrieve and delete elements in it. The function `destroy/1` will delete the database. Considering that Erlang has garbage collection, you do not need to do anything. Had the `db` module however stored everything on file, you would delete the file. We are including the `destroy` function so as to make the interface consistent. You may **not use** the `lists` library module, and have to implement all the recursive functions yourself.

Hint:

Use lists and tuples as your main data structures. When testing your program, remember that Erlang variables are single assignment. The order of the list will be implementation dependent so you may not get the same order as shown in the example.

```
db:new() -> DbRef.
db:destroy(DbRef) -> ok.
db:write(Key, Element, DbRef) -> NewDbRef.
db:delete(Key, DbRef) -> NewDbRef.
```

```
db:read(Key, DbRef) → {ok, Element} | {error, instance}.  
db:match(Element, DbRef) → [Key1, ..., KeyN].
```


Erlang Shell Session

Eshell V10.5 (abort with ^ G)

```
1> c(db).
{ok,db}
2> Db = db:new().
[]
3> Db1 = db:write(francesco, london, Db).
[{francesco,london}]
4> Db2 = db:write(lelle, stockholm, Db1).
[{francesco,london},{lelle,stockholm}]
5> db:read(francesco, Db2).
{ok,london}
6> Db3 = db:write(joern, stockholm, Db2).
[{francesco,london},{lelle,stockholm},{joern,stockholm}]
7> db:read(ola, Db3).
{error,instance}
8> db:match(stockholm, Db3).
[lelle,joern]
9> Db4 = db:delete(lelle, Db3).
[{francesco,london},{joern,stockholm}]
10> db:match(stockholm, Db4).
[joern]
11> db:write(joern, london, Db4).
[{francesco,london},{joern,london}]
```

Note:

Due to single assignment of variables in Erlang, we need to assign the up-dated database to a new variable every time.

ADVANCED: Manipulating Lists

Part A

Write a function which when given a list of integers and an integer, will return all integers smaller than or equal to that integer.

```
filter([1,2,3,4,5], 3) -> [1,2,3].
```

Part B

Write a function which given a lists will reverse the order of the elements.

```
reverse([1,2,3]) -> [3,2,1].
```

Part C

Write a function which, given a list of lists, will concatenate them.

`concatenate([[1,2,3], [], [4, five]]) → [1,2,3,4,five]`.

Hint:

You will have to use a help function and concatenate the lists in several steps.

Part D

Write a function which when given a list of nested lists, will return a flat list.

`flatten([[1,[2,[3,[]]], [[4]]], [5,6]]) → [1,2,3,4,5,6]`.

Hint:

Use the concatenate function.

ADVANCED: Database Handling using Trees

Take the `db.erl` module you wrote in exercise 2.4 and rewrite it using sorted binary trees instead of lists. Use tuples to build the trees. In this exercise we will not attempt to balance the tree to improve efficiency.

Hint:

Use the tuple `{Key, Value, LeftBranch, RightBranch}` for nodes in the tree and the atom `empty` for the leaves. The keys in `LeftBranch` are smaller than `Key` while the keys in `RightBranch` are greater than `Key`.

Note:

Make sure you save a copy of your `db.erl` module using lists somewhere else (or with a new name) before you start changing in.

3 Concurrent Programming

These exercises will help you get familiar with the syntax and semantics of concurrency in Erlang. You will solve problems that deal with spawning processes, message passing, registering, and termination. If you are having problems finding bugs or following what is going on, use the process manager.

An Echo Server

Write a server which will wait in a receive loop until a message is sent to it. Depending on the message, it should either print it and loop again or terminate. You want to hide the fact that you are dealing with a process, and access its services through a functional interface. These functions will spawn the process and send messages to it. The module `echo.erl` should export the following functions.

```
echo:start() -> ok.  
echo:stop() -> ok.  
echo:print(Term) -> ok.
```

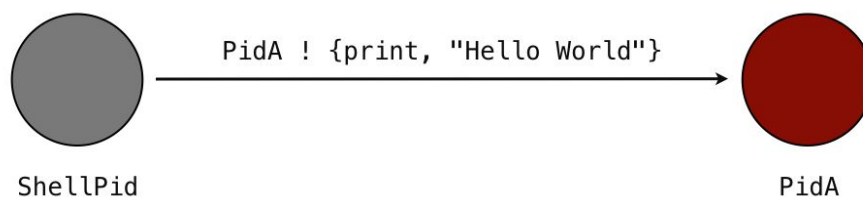


Figure 1: An Echo Server

Hint:

Use the `register/2` built in function.

Warning:

Use an internal message protocol to avoid stopping the process when you for example call the function `echo:print(stop)`.

The Process Ring

Write a program that will create N processes connected in a ring. These processes will then send M number of messages around the ring and then terminate gracefully when they receive a quit message.

Hint:

There are two basic strategies to tackling your problem. The first one is to have a central process that sets up the ring and initiates the message sending. The second strategy consists of the new process spawning the next process in the ring. With this strategy you have to find a method to connect the first process to the last.




Figure 2: The Process Ring

ADVANCED: The Process Crossing

Write a program that will create N processes connected in a ring. These processes will then send M number of messages around the ring. Halfway through the ring, however, the message will cross over the first process, which will then forward it to the second half of the ring. The ring should terminate gracefully when receiving a quit message.




Figure 3: The Process Crossing

Erlang Shell Session

Eshell V10.5 (abort with ^ G)

1> crossring:start(7,2,hello).

Process: 2 received: hello

<0.618.0>

Process: 3 received: hello

Process: 4 received: hello

Process: 1 received: hello halfway through

Process: 5 received: hello

Process: 6 received: hello

Process: 7 received: hello

Process: 1 received: hello

Process: 2 received: hello

Process: 3 received: hello

Process: 4 received: hello

Process: 1 received: hello halfway through

Process: 5 received: hello

Process: 6 received: hello

Process: 7 received: hello

Process: 1 received: hello

Process: 1 terminating

Process: 2 terminating

Process: 5 terminating

Process: 3 terminating

Process: 6 terminating

Process: 4 terminating

Process: 7 terminating

4 Process Design Patterns

These exercises will help you get familiar with process design patterns. Similar patterns will occur in different programs, and are the building blocks of systems based on OTP.

Understanding them and knowing when to use which pattern is crucial to keeping the code simple and clean.

A Database Server

Write a database server that stores a database in its loop data. You should register the server and access its services through a functional interface. Exported functions in the `my_db.erl` module should include:

```
my_db:start() -> ok.  
my_db:stop() -> ok.  
my_db:write(Key, Element) -> ok.  
my_db:delete(Key) -> ok.  
my_db:read(Key) -> {ok, Element} | {error, instance}.  
my_db:match(Element)  
-> [Key1, ..., KeyN].
```

Hint:

Use the `db.erl` module as a back end and use the server skeleton from the echo exercise.

Erlang Shell Session
<pre>Eshell V10.5 (abort with ^ G) 1> my_db:start(). ok 2> my_db:write(foo, bar). ok 3> my_db:read(baz). {error, instance} 4> my_db:read(foo). {ok, bar} 5> my_db:match(bar). [foo]</pre>

A Turnstile

Write a module, `turnstile.erl`, in which a process implements a turnstile. Model your process as a state machine with two states, *closed* and *open*. When the turnstile is in the

closed state it should accept a “coin” as input and it then transitions to the *open* state when it should accept an “enter” input when it should transition back to the *closed* state.




Figure 3: Turnstile state machine

```
turnstile:start() :: {ok, Pid}
turnstile:insert_coin(Pid, Coin) :: ok
turnstile:enter(Pid) :: ok | {error, access_denied}
```

Extend the turnstile by adding a timeout to the turnstile so if it is in the *open* state and not entered within a certain time it will automatically revert to the *closed* state.

Finally extend it to make it possible to set a certain price needed to open the turnstile. It should keep track of how much money the user has inserted before it will transition into the *open* state. If the user insert more money than needed consider implementing a way of returning the money (or keep them, that implementation detail is up to you)

ADVANCED: A Database Server with transactions

Rewrite the Database server from Exercise 4.1 to add transactions. This should be done with exported functions:

```
my_db_trans:lock() -> ok.
my_db_trans:unlock() -> ok.
```

A client starts a transaction by locking the server by calling `my_db_trans:lock` and ends it by calling `my_db_trans:unlock`. During the transaction the server will block requests from other clients and process them first after the transaction has ended.

Hint:

Use the client Process ID to pattern match inbound messages, retrieving only those from the process which has put the lock.

5 Process Error Handling

The aim of these exercises is to make you practice the simple but powerful error handling mechanisms found in Erlang. They include exiting, linking, trapping of exits and the use of catch.

The Linked Ping Pong Server

Modify the processes *A* and *B* from the file `pingpong.erl` in the exercises files given to you, by linking the processes to each other. When the `pingpong:stop` function has been called, instead of sending a *quit* message, make the first process terminate abnormally. This should result in the *exit signal* propagating to the other process, causing it to terminate as well.




Figure 6: Trapping Exits

A Reliable Mutex Semaphore

Your Mutex semaphore from the course material is unreliable. What happens if a process that currently holds the semaphore terminates prior to releasing it? Or what happens if a process waiting to execute is terminated due to an exit signal? By trapping exits and linking to the process that currently holds the semaphore, make your mutex semaphore reliable.

Hint:

When you are trapping exits, if the process has crashed before calling `link(Pid)`, the call returns true but generates a signal of the format `{'EXIT', Pid, noprocs}`.

A Reliable Database Server with transactions

Your Database server from Exercise 4.3 is unreliable. What happens if the process locking the server terminates prior to unlocking it? Use monitors to keep an eye on the locking process, making sure that the server unlocks itself and reverses any destructive operations.

ADVANCED: A Supervisor Process

Write a supervisor process that will spawn children and monitor them. If a child terminates abnormally, it will print an error message and restart it. To avoid infinite restarts (What if the Module did not exist?), put a counter which will restart a child a maximum of 5 times, and print an error message when it gives up and removes the child from its list. Stopping the supervisor should unconditionally kill all the children.

```
sup:start(SupName) - {ok, Pid}.
```

```
sup:start_child(SupName | Pid, Mod, Func, Args) - {ok, Pid}.
```

```
sup:stop(SupName | Pid) - ok.
```




Figure 7: A supervisor Process

Erlang Shell Session
<pre> Eshell V10.5 (abort with ^ G) 1> c(sup). {ok,sup} 2> sup:start(freddy). {ok,<0.41.0>} 3> {ok, Pid} = sup:start_child(freddy, my_db, init, []). {ok,<0.43.0>} 4> exit(Pid, kill). true ----- Error: Process <0.43.0> Terminated 1 time(s) Reason for termination:killed Restarting with my_db:init/0 ----- 5> {ok, Pid2} = sup:start_child(freddy, my_db, init, []). {ok,<0.47.0>} 6> i(). </pre>

Hint:

Make your supervisor start the mutex and database server processes. Note that you have to pass the function and arguments used in the spawn function, and not the start function. That might result in your process not getting registered.

If it is getting registered, kill it by using `exit(whereis(ProcName), kill)`. See if they have been restarted by calling `whereis(ProcName)` and ensuring you are getting different Process Ids every time.

If the process is not registered, kill it by calling `exit(Pid, kill)`. You will get `Pid` from the return value of the `start_child` function. (You can then start many processes of the same type). Once killed, check if the process has been restarted by calling the `i()` help function. It lists all the processes in the system, their initial function and the current function they are executing in.

Note:

`SupName` | `Pid` means you should be able to pass either the atom by which the supervisor process is registered, or the Process ID returned when the supervisor is started.

6 Functional Programming Constructs

These exercises will make you familiar with issues covered in the functional programming constructs section. In many cases, you will be using exercises created in the previous section, making the programs more robust or adding to their functionality.

Higher Order Functions

Part A

Using `fun`s and higher order functions, write a function which prints out the integers between 1 and N .

Hint:

Use `lists:seq(1, N)`.

Part B

Using `fun`s and higher order functions, write a function which, given a list of integers and an integer, will return all integers smaller than or equal to that integer.

Part C

Using `fun`s and higher order functions, write a function which prints out the even integers between 1 and N .

Hint:

Solve your problem either in two steps, or use two clauses in your `fun`.

Part D

Using `fun`s and higher order functions, write a function which, given a list of lists, will concatenate them.

Part E

Using `fun`s and higher order functions, write a function that given a list of integers returns the sum of the integers.

Hint:

Use `lists:foldl`, and try to figure out why we prefer to use `foldl` rather than `foldr`.

List Comprehensions

Part A

Using List comprehensions, create a set of integers between 1 and 10 which are divisible by three.

Example:

`[3, 6, 9]`

Part B

Using list comprehensions remove all non integers from a polymorphic list. Return the list of integers squared.

Example:

`[1, hello, 100, "boo", 9]` should return `[1, 10000, 81]`

Part C

Using list comprehensions and given two lists, return a new list which is the intersection of the two lists.

Example:

`[1, 2, 3, 4, 5]` and `[4, 5, 6, 7, 8]` should return `[4, 5]`

Hint:

Assume that the lists contain no duplicates.

Part D

Using list comprehensions and given two lists, return a new list which is the disjunction of the two lists.

Example:

`[1, 2, 3, 4, 5]` and `[4, 5, 6, 7, 8]` should return `[1, 2, 3, 6, 7, 8]`

Hint:

Assume that the lists contain no duplicates.

7 Records and Maps

Database Handling Using Records

Take the `db.erl` module you wrote in exercise 2.4 and rewrite it using records. Test it using your database server you wrote in exercise 4.1. As a record, you could use the following definition. Remember to place it in an include file.

Hint:

Use the following record definition: `-record(data, {key, value}).`

Note:

Make sure you save a copy of your `db.erl` module using lists somewhere else (or with a new name) before you start changing in.

Database Handling Using Maps

Take the `db.erl` module you wrote in exercise 2.4 and rewrite it using maps. Test it using your database server you wrote in exercise 4.1.

Note:

Make sure you save a copy of your `db.erl` module using lists somewhere else (or with a new name) before you start changing in.

8 Erlang Term Storage

Database handling using ETS

Take the `db.erl` module you wrote in exercise 2.4 and rewrite it using ETS. Test it using your database server you wrote in exercise 4.1.

Note:

Make sure that you implement exactly the same interface so you can directly use it in the database server.